

## ECE421 - Assignment 2: Neural Networks

### Part 1 - Neural Networks Using Numpy

#### 1.1 Helper Functions

1.  

```
def relu(x):  
    return np.maximum(x, 0)
```
2.  

```
#Accepts matrix with rows of o vectors  
def softmax(x):  
    #Prevent overflow by normalizing  
    x -= np.amax(x, axis=1)[:, None]  
    return np.exp(x) / np.sum(np.exp(x), axis=1)[:, None]
```
3.  

```
def computeLayer(X, W, b):  
    return W.dot(X) + b[:, None]
```
4.  

```
def CE(target, prediction):  
    #Avoid taking unnecessary logs by doing row-wise sum of the  
    matrix product  
    pred_of_tar = np.sum(target*prediction, axis=1)  
    return (-1/target.shape[0]) * np.sum(np.log(pred_of_tar))
```
5. The following is a derivation of the gradient of the cross entropy loss function with respect to the input to the softmax function  $\mathbf{o}$ . First, simplify the loss function with respect to the one non-zero term, call it the  $y^{\text{th}}$  term:

$$\begin{aligned} L &= -(1) \times \log(p_y) \\ &= -\log \frac{e^{o_y}}{\sum_{k=1}^K e^{o_k}} \end{aligned}$$

Now, the gradient must account for the partial derivatives in terms of both  $o_y$  and  $o_z, z \neq y$ :

$$\begin{aligned} \frac{dL}{do_y} &= \frac{-\sum_{k=1}^K e^{o_k}}{e^{o_y}} \times \frac{e^{o_y}(\sum_{k=1}^K e^{o_k}) - e^{2o_y}}{(\sum_{k=1}^K e^{o_k})^2} \\ &= \frac{-(\sum_{k=1}^K e^{o_k}) + e^{o_y}}{\sum_{k=1}^K e^{o_k}} \end{aligned}$$

$$\begin{aligned}
 &= -1 + \frac{e^{o_y}}{\sum_{k=1}^K e^{o_k}} \\
 \frac{dL}{do_z} &= \frac{-\sum_{k=1}^K e^{o_k}}{e^{o_y}} \times \frac{-e^{o_z+o_y}}{(\sum_{k=1}^K e^{o_k})^2} \\
 &= \frac{e^{o_z}}{\sum_{k=1}^K e^{o_k}}
 \end{aligned}$$

Noting the similarity between the two cases, the overall gradient simplifies to the following high-level expression:

$$\nabla_{\vec{o}} L = \text{softmax}(\vec{o}) - \vec{y}$$

Where  $\vec{y}$  is a vector for the one-hot encoded label. The following function computes the gradient derived:

```
def gradCE(target, prediction):
    return softmax(prediction) - target
```

## 1.2 Backpropagation Derivation

- Using the gradient of the cross entropy loss function with respect to the input to the softmax function  $\vec{o}$ , the gradient with respect to the output layer weights  $\mathbf{W}_o$  can be computed with the chain rule:

$$\begin{aligned}
 \nabla_{\mathbf{W}_o} L &= \nabla_{\vec{o}} L \cdot \frac{d\vec{o}}{d\mathbf{W}_o} \\
 &= \left( \frac{e^{\vec{o}}}{\sum_{k=1}^K e^{o_k}} - \vec{y} \right) \cdot \vec{h}^T \\
 &= (\text{softmax}(\vec{o}) - \vec{y}) \cdot \vec{h}^T
 \end{aligned}$$

Note that this gradient has the form of a K by H matrix, or equivalently can be thought of as a set of K number of H sized vectors.

- Similarly to the derivation above, the following is a derivation of the cross entropy loss function with respect to the output layer biases:

$$\begin{aligned}
 \nabla_{\mathbf{b}_o} L &= \nabla_{\vec{o}} L \cdot \frac{d\vec{o}}{d\mathbf{b}_o} \\
 &= \left( \frac{e^{\vec{o}}}{\sum_{k=1}^K e^{o_k}} - \vec{y} \right) \circ \vec{1} \\
 &= \text{softmax}(\vec{o}) - \vec{y}
 \end{aligned}$$

3. In order to find the gradient of the cross entropy loss function with respect to the hidden layer weights  $\mathbf{W}_h$ , the loss was expressed in terms of the hidden layer weights  $\mathbf{W}_h$  and the derivative taken directly. The one non-zero loss term has subscript  $y$ .

$$L = -\log\left(\frac{N}{D}\right)$$

$$= -\log\left(\frac{\exp(W_y^{(o)} \cdot \text{ReLU}(W^{(h)} \bar{x} + b^{(h)}) + b_y^{(o)})}{\sum_{k=1}^K \exp(W_k^{(o)} \cdot \text{ReLU}(W^{(h)} \bar{x} + b^{(h)}) + b_k^{(o)})}\right)$$

Then:

$$\frac{dL}{dW_{ij}^{(h)}} = -\frac{D}{N} \times \left[ \frac{N \times \text{ReLU}'(W^{(h)} \bar{x} + b^{(h)}) \cdot (x_j)(W_{yi}^{(o)}) \times D}{D^2} - \frac{D \times \text{ReLU}'(W^{(h)} \bar{x} + b^{(h)}) \cdot (x_j)(W_{ki}^{(o)}) \times N}{D^2} \right]$$

$$= -\text{ReLU}'(\bar{h}) \cdot \left( \frac{x_j W_{yi}^{(o)} \sum_{l=1}^K \exp(o_l)}{\sum_{k=1}^K \exp(o_k)} - \frac{\sum_{l=1}^K \exp(o_l) x_j W_{li}^{(o)}}{\sum_{k=1}^K \exp(o_k)} \right)$$

$$= \text{ReLU}'(\bar{h}) \cdot x_j \left( \sum_{l=1}^K \frac{\exp(o_l) W_{li}^{(o)}}{\sum_{k=1}^K \exp(o_k)} - W_{yi}^{(o)} \right)$$

Here is the equivalent high-level vectorized expression:

$$\text{ReLU}'(\bar{h}) \circ \left[ \left( (W^{(o)})^T \cdot \text{softmax}(\bar{o}) - (W_y^{(o)})^T \right) \cdot \bar{x}^T \right]$$

4. Similarly to the derivation above, The following expression is the gradient of the cross entropy loss function with respect to the hidden layer biases.

$$\frac{dL}{db_i^{(h)}} = \text{ReLU}'(\bar{h}) \cdot \left( \sum_{l=1}^K \frac{\exp(o_l) W_{li}^{(o)}}{\sum_{k=1}^K \exp(o_k)} - W_{yi}^{(o)} \right)$$

This is the equivalent high-level vectorized expression:

$$\text{ReLU}'(\bar{h}) \circ \left[ (W^{(o)})^T \cdot \text{softmax}(\bar{o}) - (W_y^{(o)})^T \right]$$

### 1.3 Learning

The following (Python) function was used for training the neural network. (Runtime for 200 epochs was roughly 5 minutes).

```
def gradDescent():
    #Constants
    H = 1000
    K = 10
    F = 784
```

```
EPOCHS = 200
gamma = 0.9
alpha = 0.1
N = 10000

#Training and validation data/target extraction
Data = loadData()
data = Data[0].reshape(N, 784)
validData = Data[1].reshape(6000, 784)
Targets = convertOneHot(Data[3], Data[4], Data[5])
target = Targets[0]
validTarget = Targets[1]
del Data
del Targets

#Initialization of weights and biases
Wo = np.random.normal(0, np.sqrt(2/H), (K, H))
Wh = np.random.normal(0, np.sqrt(2/F), (H, F))
bo = np.zeros(K)
bh = np.zeros(H)

#Initialization of V matrices
VWo = np.full((K, H), 1e-5)
VWh = np.full((H, F), 1e-5)
Vbo = np.full(K, 1e-5)
Vbh = np.full(H, 1e-5)

#Data containers used for plotting
xpoints = np.arange(1, EPOCHS+1)
ytrain = []
yvalid = []

#Local function for computing a forward pass.
#Training is a boolean (set false for validation).
def forward_propagation(training):
    nonlocal data, Wh, bh, Wo, bo
    if training:
        Sh = computeLayer(data.T, Wh, bh)
    else:
        Sh = computeLayer(validData.T, Wh, bh)
```

```
Xh = relu(Sh)
So = computeLayer(Xh, Wo, bo)
Xo = softmax(So.T)

return Sh, Xh, So, Xo

#Main training loop
for epoch in range(EPOCHS):
    #Forward propagation
    Sh, Xh, So, Xo = forward_propagation(True)

    #Gradients/backpropagation
    grad_bo = gradCE(target, So.T)
    grad_Wo = grad_bo.T.dot(Xh.T)/N
    grad_bo = np.sum(grad_bo, axis=0)/N
    grad_bh = (Wo.T.dot(Xo.T) - Wo.T.dot(target.T))*np.array(Xh,
dtype=bool)
    grad_Wh = grad_bh.dot(data)/N
    grad_bh = np.sum(grad_bh, axis=1)/N

    #Gradient descent with momentum
    VWo = gamma*VWo + alpha*grad_Wo
    VWh = gamma*VWh + alpha*grad_Wh
    Vbo = gamma*Vbo + alpha*grad_bo
    Vbh = gamma*Vbh + alpha*grad_bh

    #Update Weights and biases
    Wo -= VWo
    Wh -= VWh
    bo -= Vbo
    bh -= Vbh

    #Add new plotting data
    ytrain.append(accuracy(target, Xo))
    yvalid.append(accuracy(validTarget,
forward_propagation(False)[3]))

#Plot findings. Plots shown/labeled in main() function
plt.plot(xpoints, ytrain, label = "Training")
plt.plot(xpoints, yvalid, label = "Validation")
```

The following function was used for computing the accuracy for a given set of predictions and targets:

```
def accuracy(target, prediction):  
    #N = number of examples  
    N = target.shape[0]  
    return np.sum(target[np.arange(N), np.argmax(prediction, axis=1)]) / N
```

The accuracy and loss was plotted against epochs for both training and validation. The hidden unit size (H) was set to 1000, momentum parameter ( $\gamma$ ) was set to 0.9, and learning rate ( $\alpha$ ) was set to 0.1.

